

Week1day2

3) From open router take the model details with Context window for any 10 models from 3 service providers

1)GPT – 6 Astra

Open AI , 1.05M context

2)GPT – 6 Astra Pro

Open AI,1.05M context

3) MAI – Image 2.6 Flash

Microsoft 4k context

4)Ling 3.0 Flash Sante [free]

InclusionAI 262kcontext

5) Qwen3.8 Max(0902)

QWEN 1M context

6)Muse Spark1.3 Contributor

Meta 1.05M context

7)Gemini 3.8 Flash

Google 1.05M context

8)Claude Fable 5.1(batch)

Anthopic 1M context

9)Mercury 2.5Preview

inception 260kcontext

10)Granite 4.2 8B

IBM-Granite 131kcontext

2) Hallucination check for the below prompt

[Write a 4-line poem where every line ends with a word that

strictly rhymes with the word 'Silver'.]

Deepseek response

chat.deepseek.com/a/chat/s/6856a927-7521-4923-a05f-1e7855d865c9

deepseek

New chat

Open Weight LLM List

Instant

Write a 4-line poem where every line ends with a word that strictly rhymes with the word 'Silver'

While "silver" is famously one of the few English words with no perfect rhymes, here is a 4-line poem using **near rhymes** (assonance) that echo its sound:

The moonlight gleams on calm **water**,
The stream reflects the stars that **quiver**,
A whispered dream, the night grows **softer**,
The world remains a gift to **deliver**.

Message DeepSeek

DeepThink Smart Search

AI-generated, for reference only

QWEN Response

chat.qwen.ai/c/guest

Qwen3.7-Plus

Write a 4-line poem where every line ends with a word that strictly rhymes with the word 'Silver'.

Thinking completed

The autumn wind sweeps down the winding river,
And makes the golden leaves begin to shiver.
I watch the trembling aspen branches quiver,
Awaiting dawn and what the light will deliver.

OPEN AI- GPT

https://chatgpt.com/c/6a9d7e2b-6784-83e9-920b-327a75ab2328

PT


chat

Free offer

Write a 4-line poem where every line ends with a word that strictly rhymes with the word 'Silver'.

Moonlight shimmers softly on the **silver**,
Stars dance above, growing ever **chillier**,
Dreams drift through the night, becoming **still-er**,
Dawn paints the sky with a glow ever **brighter**.

Gemini response



1) Get the List of Open weight LLM from different service providers

QWEN 3.7Plus response

Here is a comprehensive list of **open-weight Large Language Models (LLMs)** grouped by their primary service providers, reflecting the landscape as of 2025–2026.

***Note:** "Open-weight" means the model's trained parameters (weights) are publicly available for download, local deployment, and fine-tuning, even if the full training data or code is not entirely open source. Licensing terms vary significantly between providers.*

1. **Meta**

Meta has been a major champion of open-weight models, providing highly competitive backbones for the community.

Llama 3 / Llama 4 Series: Includes Llama 3 (8B, 70B) and the newer Llama 4 (including Llama 4 Scout), which are designed for scalability, efficiency, and natively multimodal capabilities [[1]], [[27]], [[41]].

Muse Series: Includes models like Muse Glimmer (e.g., 30B parameters), designed specifically to run efficiently on local hardware [[268]].

License: Meta Llama Community License (free for most uses, but restricts usage for companies with >700M monthly active users and prohibits using outputs to train other LLMs) [[26]].

2. **Alibaba (Qwen)**

Alibaba's Qwen (Tongyi Qianwen) series has become a dominant force in the open-weight space, frequently topping leaderboards.

Qwen Series: Ranging from Qwen1.5 to the latest Qwen3.x and Qwen3.8 models (e.g., Qwen3.8 Max) [[2]], [[43]].

Strengths: Highly regarded for advanced reasoning, coding capabilities, and switchable "thinking" modes, often matching or beating top closed-source models [[21]], [[43]].

License: Custom Qwen License (free for commercial use under 100M monthly active users) [[14]].

3. **Mistral AI**

The European AI leader is known for highly efficient, dense, and Mixture-of-Experts (MoE) architectures.

Mistral & Mixtral Series: Includes the foundational Mistral 7B, Mixtral 8x7B, and Mixtral 8x22B (MoE) [[14]].

Newer Releases: Mistral Large 3, Mistral Medium 3.5, and Mistral Small 4, which continue to push the boundary of high-performance, compute-efficient open-weight models [[4]], [[31]], [[192]].

License: Predominantly Apache 2.0, making it highly permissive for commercial use [[14]].

4. **Microsoft**

Microsoft focuses on "Small Language Models" (SLMs) that punch above their weight class in reasoning.

Phi Series: Includes Phi-2, Phi-3 (Mini, Small, Medium), and Phi-4 [[14]], [[41]].

Strengths: Optimized for high reasoning and complex task performance with remarkably few parameters, making them ideal for edge devices, laptops, and local deployment [[35]].

License: MIT License (highly permissive) [[14]].

5. **Google**

Google provides lightweight, state-of-the-art open-weight models built from the same research and technology used to create their proprietary Gemini models.

Gemma Series: Includes Gemma, Gemma 2, and the recently announced Gemma 4 [[198]].

Strengths: Efficient, transparent, and designed for both research and commercial deployment. Gemma 4 notably switched to a more permissive Apache 2.0 license [[198]], [[199]].

6. **DeepSeek**

DeepSeek has disrupted the market by offering frontier-level performance at a fraction of the computational cost.

- * **DeepSeek Series**:

Includes DeepSeek V2, V3, and V4 (along with specialized variants like DeepSeek V4 Pro and Flash) [[32]], [[41]].

- * **Strengths**:

Highly efficient Mixture-of-Experts (MoE) architecture, excelling in advanced reasoning, mathematics, and coding [[41]].

- * **License**:

Custom license (free for commercial use, but derivatives must be disclosed and remain subject to the license) [[14]].

7. **OpenAI**

Historically closed, OpenAI has recently introduced a line of open-weight models for specific deployment scenarios.

- * **GPT-OSS Series**:

Includes `gpt-oss-20B` (designed to run on high-end mobile devices) and `gpt-oss-120B` (built for production general-purpose reasoning) [[29]], [[35]].

- * **License**:

Standard Apache 2.0 license, marking a significant shift in their model distribution strategy [[11]].

8. **Moonshot AI**

- * **Kimi Series**:

Includes Kimi-K2.6, Kimi-K2.7-Code, and Kimi-K3 [[5]], [[30]], [[231]].

- * **Strengths**:

Positioned as long-context, agent-oriented LLMs specifically optimized for complex, multi-step coding tasks and autonomous multi-agent swarms [[206]], [[208]].

9. **Z.ai (Zhipu AI)**

- * **GLM Series**:

Includes GLM-4, GLM-5.1, and GLM-5.2 [[200]], [[234]].

- * **Strengths**:

Built for long-horizon tasks and agentic workflows, frequently ranking at the top of open-weight coding and reasoning benchmarks (e.g., SWE-Bench) [[200]], [[236]].

10. **NVIDIA**

- * **Nemotron Series**:

Includes Nemotron 3 Ultra [[228]].

- * **Strengths**:

Optimized for enterprise AI workloads, tool use, and synthetic data generation. Often features innovative quantization (e.g., NVFP4) for highly efficient deployment [[228]], [[229]].

11. **Other Notable Providers**

- * **Apple**:

OpenELM, an efficient language model family with open weights and training code, using a layer-wise scaling strategy designed to run efficiently on-device (e.g., iPhones, Macs) [[44]], [[46]].

* **Xiaomi**: **MiMo Series** (e.g., MiMo-V2 Pro), featuring up to 1 trillion parameters and massive 1M+ token context windows, released under an MIT license [[195]], [[214]].

* **Tencent**: **Hy Series** (e.g., Hy3, Hy4-preview), Apache-licensed models competing in the open-weight space with strong general and reasoning performance [[239]], [[271]].

* **Upstage**: **Solar Series** (e.g., Solar-Open2-250B), open-weight models built for autonomous tasks and enterprise use [[258]], [[259]].

* **Poolside**: **Laguna S**, a specialized open-weight coding model designed to outperform rivals significantly larger than its size [[9]], [[254]].

🔑 Key Considerations When Choosing an Open-Weight LLM

1. **Licensing**: Always verify the license. "Open-weight" does not always mean "Open Source" (OSI-defined). Models like **Mixtral**, **Gemma 4**, and **Phi** use permissive licenses (Apache 2.0 / MIT), while **Llama** and **Qwen** use custom community licenses with specific commercial or derivative-training restrictions [[14]], [[26]].
2. **Hardware Requirements**: Models range from ~2B–3B parameters (runnable on consumer laptops or high-end phones) to 100B+ parameters (requiring multi-GPU enterprise setups or heavy quantization) [[35]].
3. **Specialization**: Choose based on your use case. For example, use **Phi** for edge reasoning, **Kimi/GLM** for long-context coding agents, **Nemotron** for synthetic data pipelines, and **Llama/Qwen** for general-purpose, highly fine-tunable backbones.

For live performance comparisons, you can refer to community-maintained leaderboards like the [Hugging Face Open LLM Leaderboard](https://huggingface.co/spaces/open-llm-leaderboard/open-llm-leaderboard) or [Vellum Open LLM Leaderboard](https://www.vellum.ai/open-llm-leaderboard) [[3]], [[32]].

Deepseek response

Here is a list of notable open-weight LLMs from various providers, including their key specs and licensing information.

Featured Open-Weight LLMs

Model	Provider	Key Specs	License	Notes
Inkling	Thinking Machines Lab	975B total / 41B active params (MoE); 1M context; Multimodal	Apache 2.0	New open-weight model aimed at enterprise customization and fine-tuning.
Gemma 4		31B, 26B (MoE); E2B, E4B for edge	Apache 2.0	Strong reasoning performance. Compact versions for edge devices.
Qwen3	Alibaba	235B, 32B, 14B, 8B, etc.	Apache 2.0	Strong general performance across sizes.
Llama 3.3 / 4	Meta	70B, 109B (MoE - Scout), 400B (Maverick)	Llama Community License	Very popular; has a commercial use clause for large deployers.
DeepSeek	DeepSeek	R1 (67B), V3.2 (~685B MoE); 1M context	MIT	Known for strong reasoning and coding performance.
Mistral	Mistral AI	Small 3.1 (24B), Large (123B), Ministral (8B)	Apache 2.0	High-performance models with a focus on efficiency.
Phi-4	Microsoft	14B, mini (3.8B)	MIT	Compact models with strong mathematical and reasoning capabilities.
EuroLLM	European Consortium	1.7B, 9B, 22B	Fully Open	Designed for strong performance across all 24 official EU languages.
Apertus	EPFL / ETH Zurich	8B, 70B	Apache 2.0	Fully open model (architecture, weights, training data) supporting over 1000 languages.
GLM	Zhipu AI	4.5-Air (12B), 5.1 (~754B MoE)	MIT	Strong Chinese and English language performance.
gpt-oss	OpenAI	20B, 120B	Apache 2.0	OpenAI's first open-weight model; high reasoning benchmarks.

Key Considerations When Choosing

- * **License**: This determines if and how you can use a model commercially. **Apache 2.0** and **MIT** are generally the safest for unrestricted commercial use. Models like Llama come with community licenses that may have user-count clauses.
- * **Hardware Requirements**: Running these models requires significant compute. You can check a model's VRAM needs to ensure it fits your infrastructure, often using quantization (e.g., Q4_K_M) to reduce requirements.
- * **Hosting**: You can run these models locally using tools like **Ollama** or **llama.cpp**, or deploy them in production with frameworks like **vLLM** or **SGLang**.

OPEN AI response

Yes. If by **“open-weight LLM”** you mean models whose trained weights can be downloaded and run/hosted by users (with licensing varying by provider), there are now many strong options across US, Chinese, European, and other providers. Hugging Face

currently lists thousands of downloadable/inference-ready models, including Llama, Qwen,



DeepSeek, Gemma, Mistral and others. [Hugging Face](#)

Major Open-Weight LLM Providers

Provider / Organization	Country / Region	Major open-weight model families	Typical sizes	Best known for
Meta	US	Llama 3.x, Llama 4	1B–400B+	General-purpose, ecosystem
Alibaba / Qwen	CN China	Qwen 2.5, Qwen 3, Qwen-Coder, Qwen-VL	0.5B–400B+	General AI, coding, multilingual
DeepSeek	CN China	DeepSeek V3, V3.2, R1, Coder	7B–670B+	Reasoning, coding, MoE
Mistral AI	FR France	Mistral, Mixtral, Mistral Small/Large	7B–675B	Efficient models, European AI
Google	US	Gemma 2, Gemma 3, Gemma 4	1B–120B+	Small/edge, multimodal
Microsoft	US	Phi-3, Phi-4	~3B–15B	Small efficient models
IBM	US	Granite	~2B–34B+	Enterprise, coding, business
NVIDIA	US	Nemotron	Various	Reasoning, enterprise, agentic AI

Provider / Organization	Country / Region	Major open-weight model families	Typical sizes	Best known for
01.AI	CN China	Yi	6B–34B	General-purpose multilingual
Zhipu AI	CN China	GLM / GLM-4.x/5.x	Various	Reasoning, agents
Moonshot AI	CN China	Kimi K2/K2.x	Large MoE	Reasoning, coding, agents
Tsinghua / THU DM	CN China	ChatGLM / GLM	Various	Chinese + multilingual
01.AI / community	CN China	Yi series	6B–34B	General LLM
Cohere	CA Canada	Command R / R+	~35B–104B	RAG, enterprise
Databricks	us US	DBRX	132B	MoE, enterprise
AI2 (Allen Institute)	us US	OLMo	1B–32B	Truly open research
Technology Innovation Institute	AE UAE	Falcon	7B–180B	Open research/LLMs
AI21 Labs	IL Israel	Jamba	Various	Hybrid SSM/Transformer
Snowflake	us US	Arctic	480B MoE	Enterprise MoE
Upstage	KR South Korea	Solar	10B–22B	Efficient enterprise models
Sakana AI	JP Japan	Various open models	Various	Model merging/research
Baidu	CN China	ERNIE / ERNIE 4.x	Various	Chinese language/AI

Provider / Organization	Country / Region	Major open-weight model families	Typical sizes	Best known for
ByteDance	CN China	Seed / Doubao-related models	Various	Multimodal/agents
Tencent	CN China	Hunyuan	Various	Multimodal/general AI
MiniMax	CN China	MiniMax models	Large MoE	General/agent AI

Particularly important families

1. Meta — Llama

- Llama 3.1
- Llama 3.2
- Llama 3.3
- Llama 4 Scout
- Llama 4 Maverick

Llama remains one of the largest open-weight ecosystems, with extensive support from inference frameworks and cloud providers. Hugging Face currently shows Llama 3.1 8B,



Llama 3.2 1B/3B and Llama 3.3 70B among heavily used models. Hugging Face

2. Alibaba — Qwen

- Qwen 2.5
- Qwen 2.5 Coder
- Qwen 2.5-VL
- Qwen 3
- Qwen 3 Coder
- Qwen 3.5

Qwen is particularly interesting if you need **coding + multilingual + reasoning**. Qwen 3 models from sub-1B through larger models are currently prominent on Hugging Face. **H**





Hugging Face

3. DeepSeek

- DeepSeek-V3
- DeepSeek-V3.x
- DeepSeek-R1
- DeepSeek-R1-Distill
- DeepSeek-Coder

DeepSeek is particularly important for **reasoning and MoE architectures**. Its large models can have hundreds of billions of total parameters while activating substantially fewer parameters per token, which makes inference more economical than the raw parameter



count suggests. I  IT Pro

4. Mistral AI

- Mistral 7B
- Mixtral
- Mistral Small
- Mistral Large
- Mistral 3
- Codestral / Devstral-related models

Mistral is one of the most important European open-weight model providers. Its models are also widely supported by open-source inference stacks.

5. Google — Gemma

- Gemma
- Gemma 2
- Gemma 3
- Gemma 4
- CodeGemma
- PaliGemma

Gemma is particularly attractive when you need **small models that can run locally**, including on relatively modest GPUs/edge hardware. Current Hugging Face listings include Gemma



Hugging Face

models from 1B upward. **H**

6. Microsoft — Phi

- Phi-3 Mini
- Phi-3 Small
- Phi-3 Medium
- Phi-4
- Phi-4-mini

Good choice for **small-footprint local applications**, especially when a 7B–70B model would be unnecessarily expensive.

7. IBM — Granite

Granite is aimed heavily at **enterprise workloads, coding, RAG and business applications**.

8. AI2 — OLMo

OLMo is especially noteworthy if your definition of “open” is stricter than simply downloadable weights. AI2 has emphasized openness around the research/data/training ecosystem rather than only releasing model weights.

Open-weight models worth evaluating

If you're building a shortlist today, I'd start with:

Use case	Models to evaluate
 General-purpose	Qwen 3/3.5, Llama 4, Mistral
 Reasoning	DeepSeek R1/V3.x, Qwen reasoning models, GLM
 Coding	Qwen Coder, DeepSeek Coder, Devstral
 Small/local	Gemma 3/4, Phi-4, Qwen small models
 Vision	Qwen-VL, Gemma, Llama 4, Mistral multimodal models
 Enterprise	Granite, Llama, Mistral, Qwen
 Multilingual	Qwen, Llama, Gemma, Mistral
 Research/open science	OLMo, Mistral, Llama, Qwen
CN Chinese	Qwen, DeepSeek, GLM, Kimi, Hunyuan
EU European	Mistral

Use case

Models to evaluate

us US

Llama, Gemma, Phi, Granite, OLMo

One important distinction: **“open-weight” does not necessarily mean “open-source.”** A model may make its weights downloadable while imposing restrictions through its license.

Before commercial deployment, check the specific model's license rather than assuming all downloadable models have the same rights.

For example, Hugging Face's current supported-model ecosystem includes DeepSeek, Llama, Gemma, Mistral, Qwen, Phi and many others, illustrating how broad the deployable open-



weight ecosystem has become. **H**

Hugging Face

If you're doing this for **enterprise/vendor evaluation**, I can also produce a much more useful **2026 Open-Weight LLM comparison with ~30–50 models**, including **provider, model, parameters, active parameters, context window, license, benchmark scores, GPU requirements, inference cost, coding/reasoning capability, and Hugging Face/GitHub links**.